

29TH USENIX
SECURITY SYMPOSIUM




ParmeSan: Sanitizer-guided Greybox Fuzzing


Sebastian Österlund, Kaveh Razavi, Herbert Bos, Cristiano Giuffrida
Vrije Universiteit Amsterdam

Slides made by Manh-Dung Nguyen


Coverage-guided Greybox Fuzzing




Directed Greybox Fuzzing




Coverage-guided Greybox Fuzzing



Directed Greybox Fuzzing



Sanitizer-guided Greybox Fuzzing

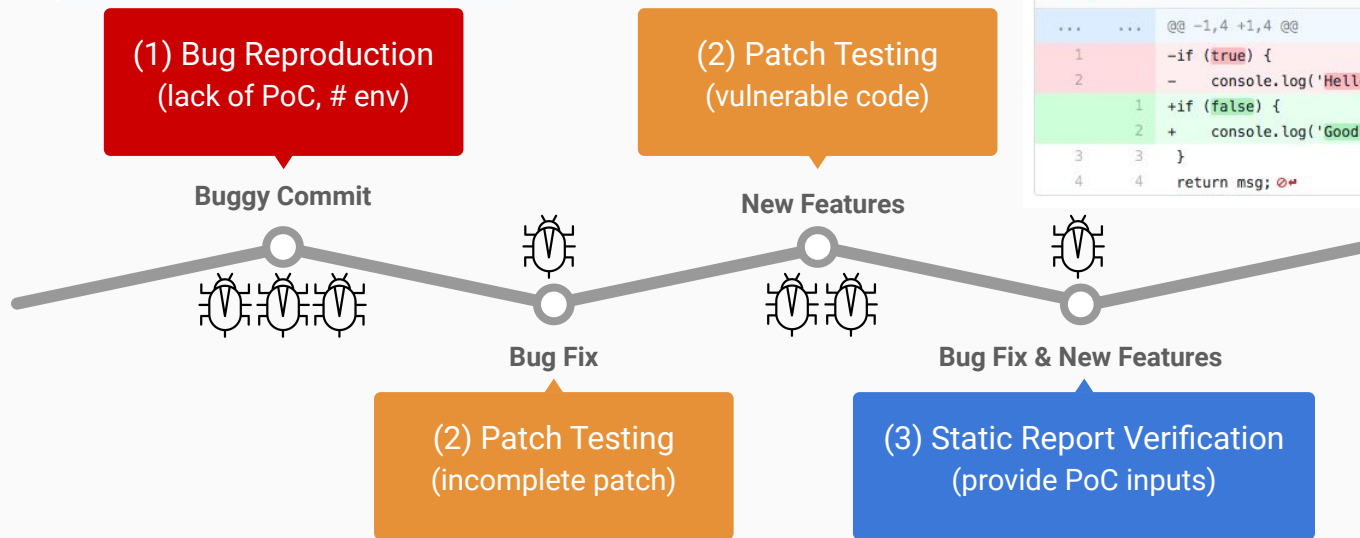


*Better view
of bugs*

*Better guidance
towards bugs*

Directed Greybox Fuzzing (DGF)

🦋 CVE-2018-20623 Detail



```
4 nodes/ba780c3a.39e07/func
... @@ -1,4 +1,4 @@
1 -if (true) {
2 - console.log('Hello world.')
1 +if (false) {
2 + console.log('Goodbye world.')
3 }
4 return msg; 🦋
```

Focus on (3) Static Report Verification

A "NullPointerException" could be thrown; "providedClass" is nullable here.

		Reliability	Security	Maintainability	New code Since last release
🦋 Bugs	2	B	1	B	
🔒 Security Vulnerabilities	0	A	0	A	
🔥 Security Hotspots	39	-	0	-	
🕒 Technical Debt	6 days	C	0	A	
🐛 Code Smells	319	-	0	-	

Goal & Challenges



Sanitizer-guided Greybox Fuzzing

- Efficient interesting targets identification
- Efficient bug detection
- Low instrumentation-/runtime- overhead



Challenges

- No ideal sanitizers
 - Specific sanitizer: focus on specific bug & fewer targets
 - General sanitizer: more bugs & more targets
- More targets (compared to bug reproduction & patch testing)

Contributions



Techniques

- ★ A generic way of *finding interesting fuzzing targets* by relying on existing compiler sanitizer passes.
- ★ A *dynamic* approach to build a *precise* control-flow graph used to steer the input towards our targets.

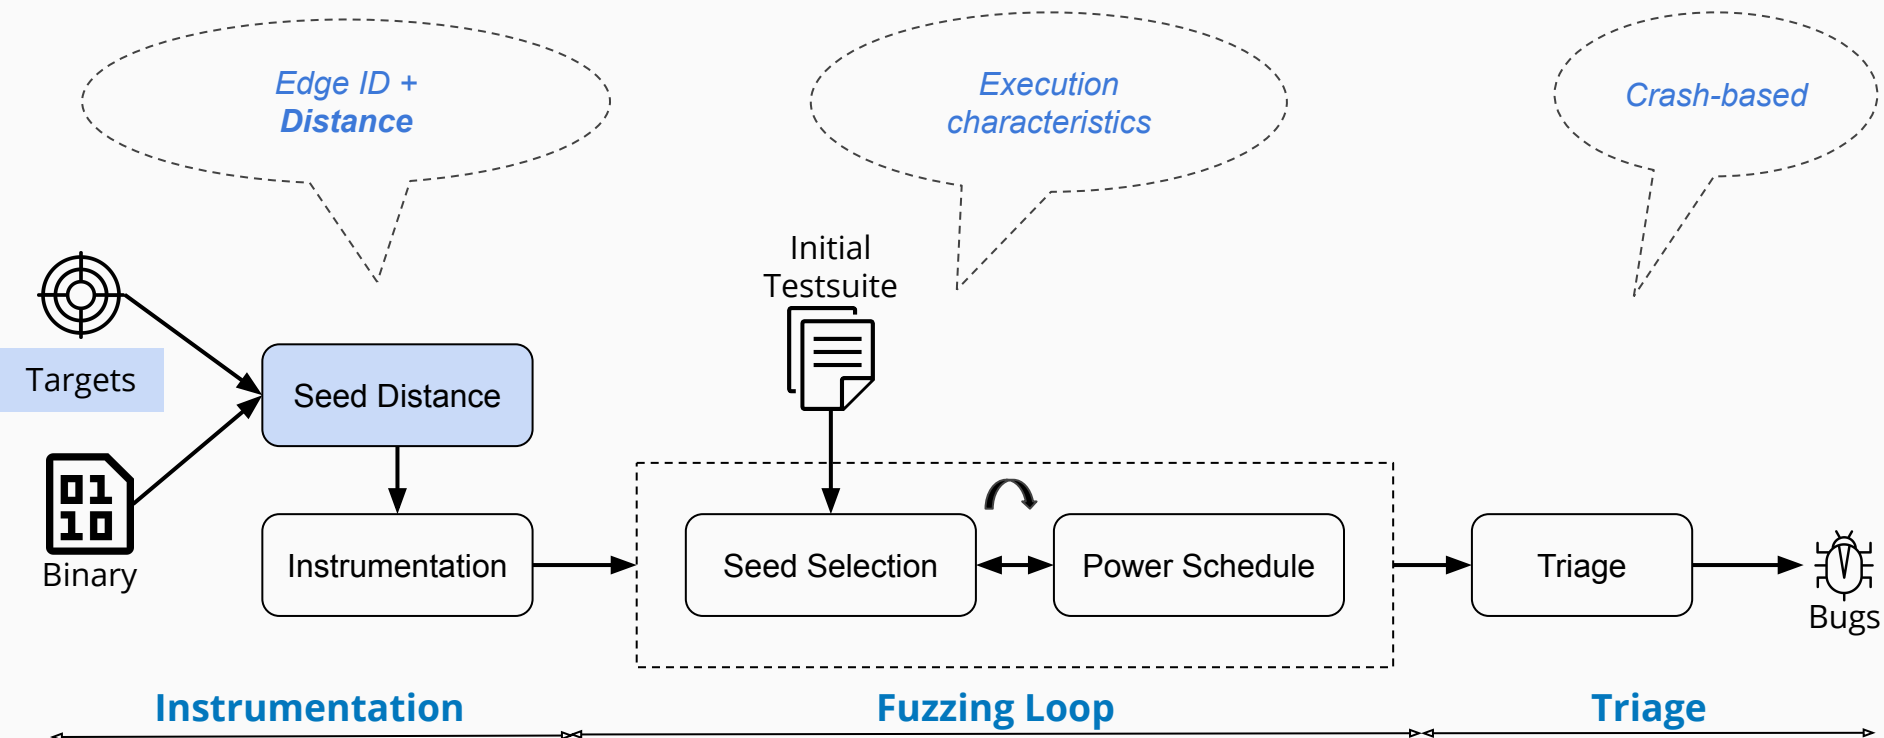
Implementation

- ★ ParmeSan, the *first sanitizer-guided fuzzer* using a two-stage directed fuzzing strategy to efficiently reach all the interesting targets.

Evaluation

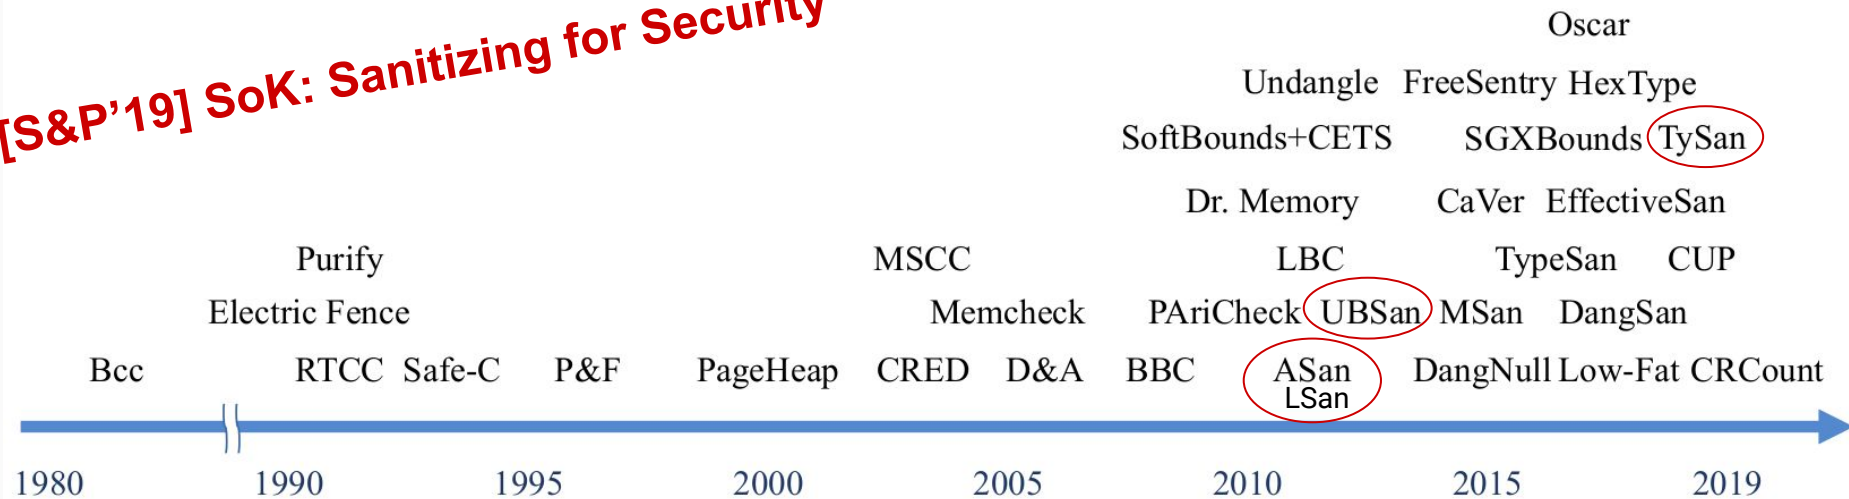
- ★ ParmeSan finds the *same bugs* as state-of-the-art coverage-guided and directed fuzzers *in less time*.

Background: Directed Greybox Fuzzing




Background: Sanitizers - Dynamic Analysis Tools (1)


[S&P'19] SoK: Sanitizing for Security



C/C++ Bugs



UndefinedBehaviorSanitizer (UBSan)

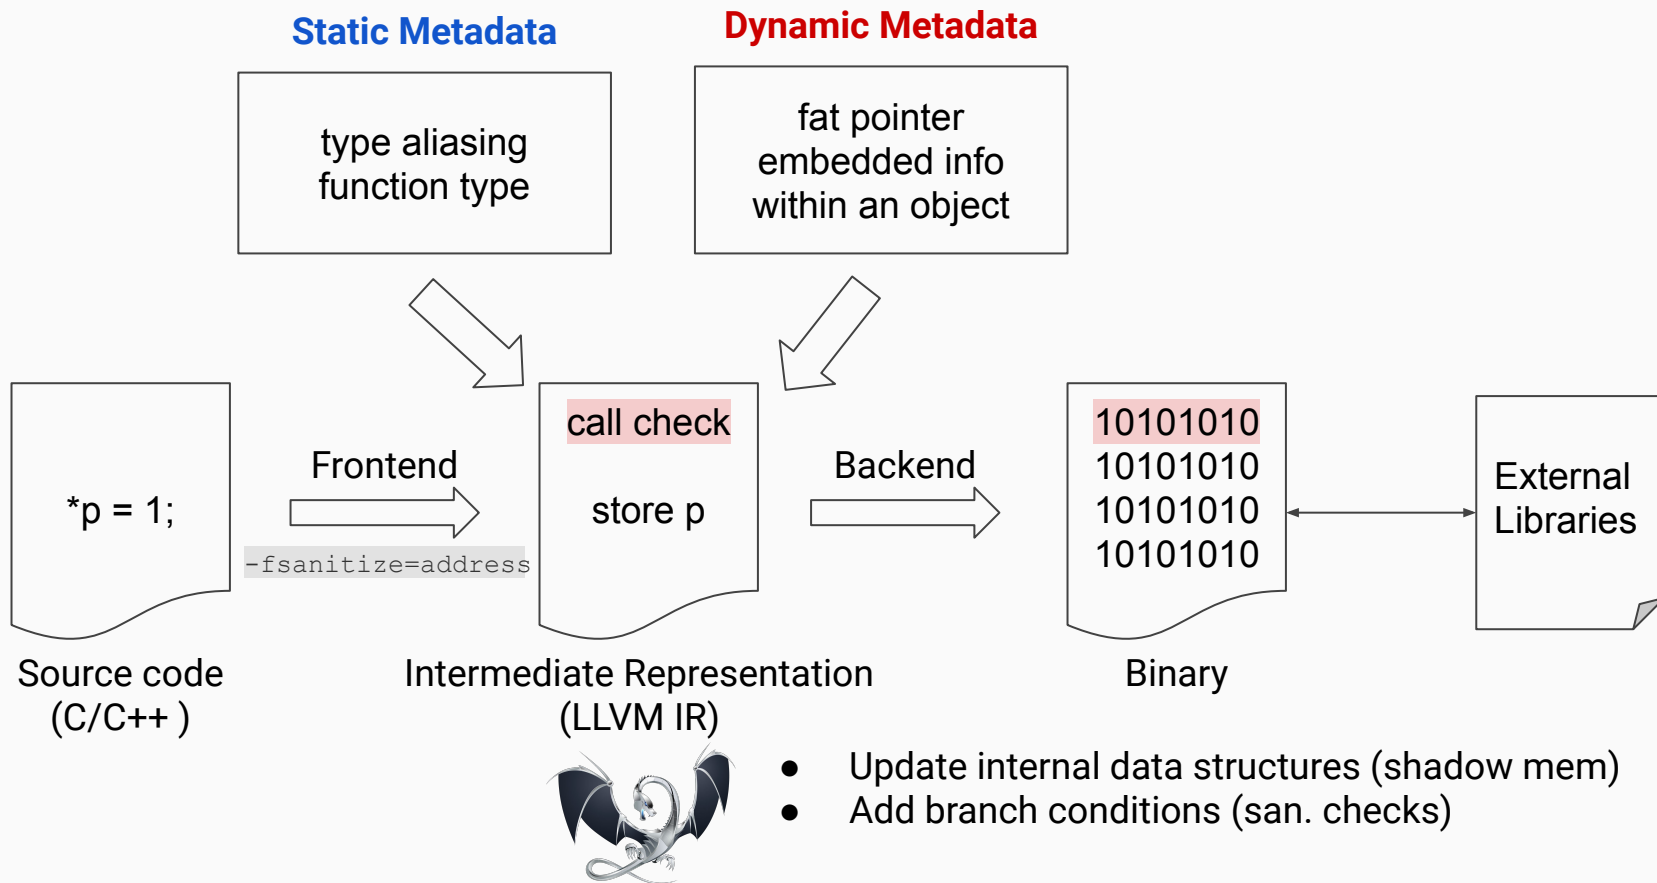


AddressSanitizer (ASan)
LeakSanitizer (LSan)

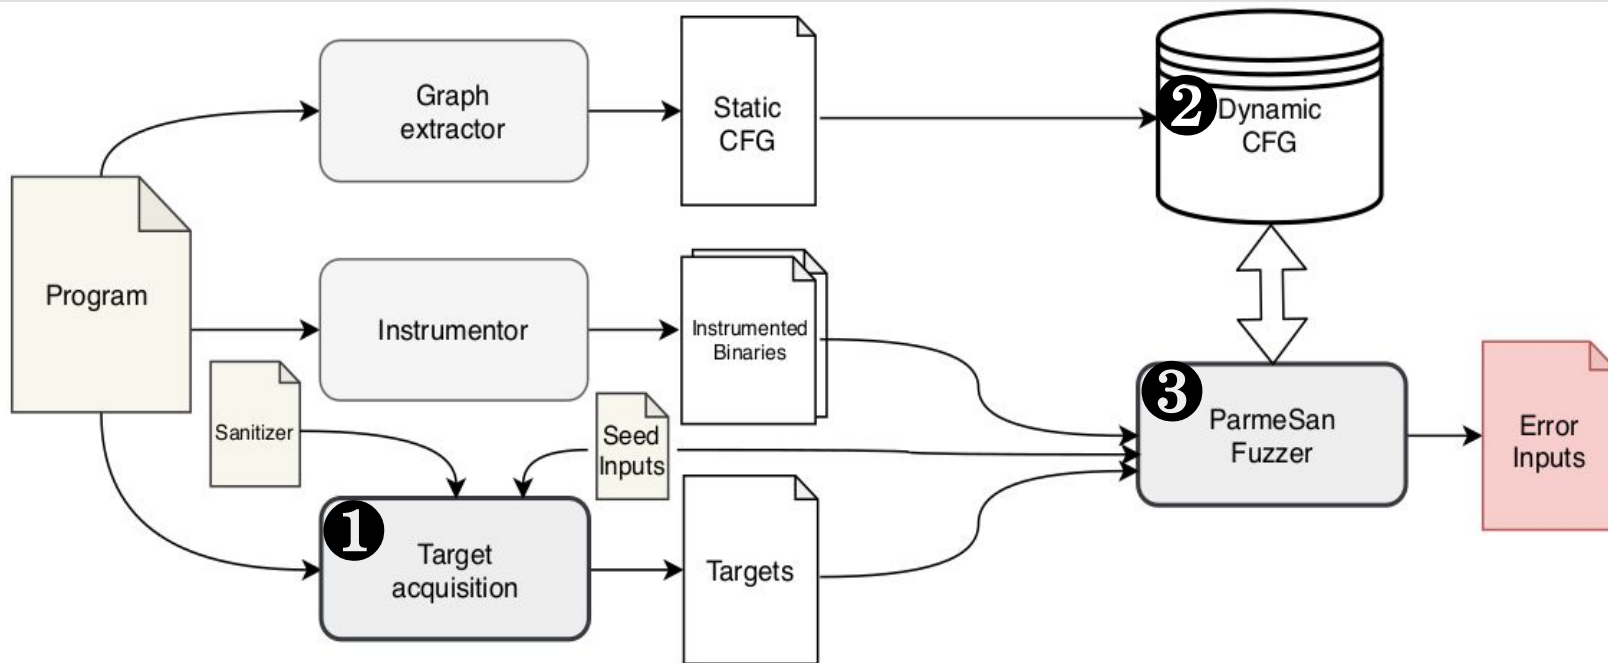


TypeSanitizer (TySan)

Background: Sanitizers - Instrumentation & Metadata (2)

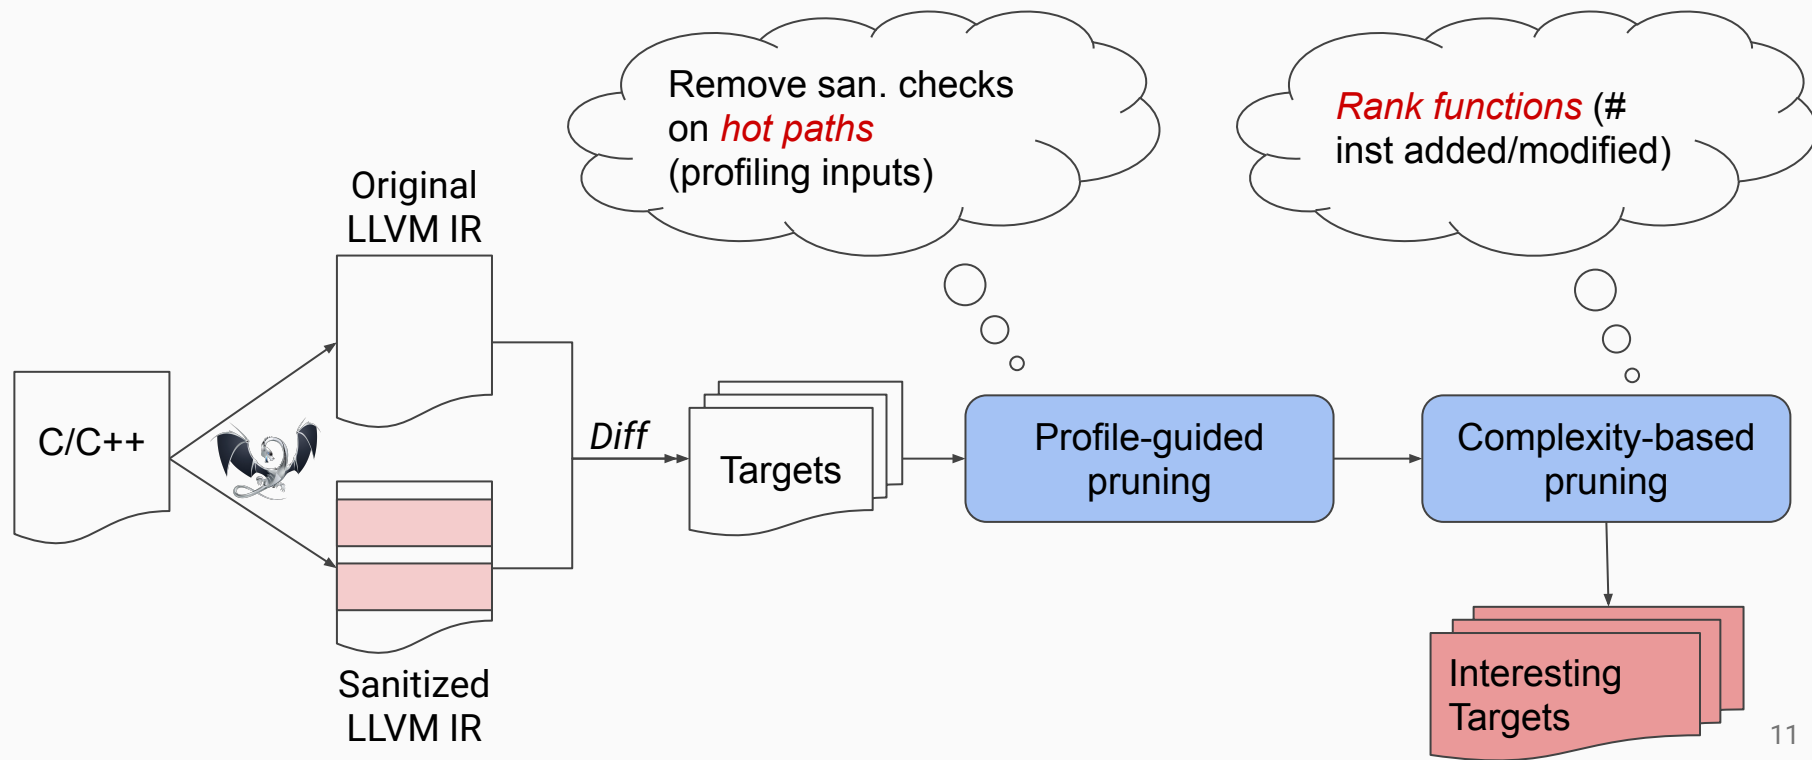


ParmeSan Overview



Target Acquisition

Goal: Find *interesting (likely-buggy)* targets (BBs) for directed fuzzing

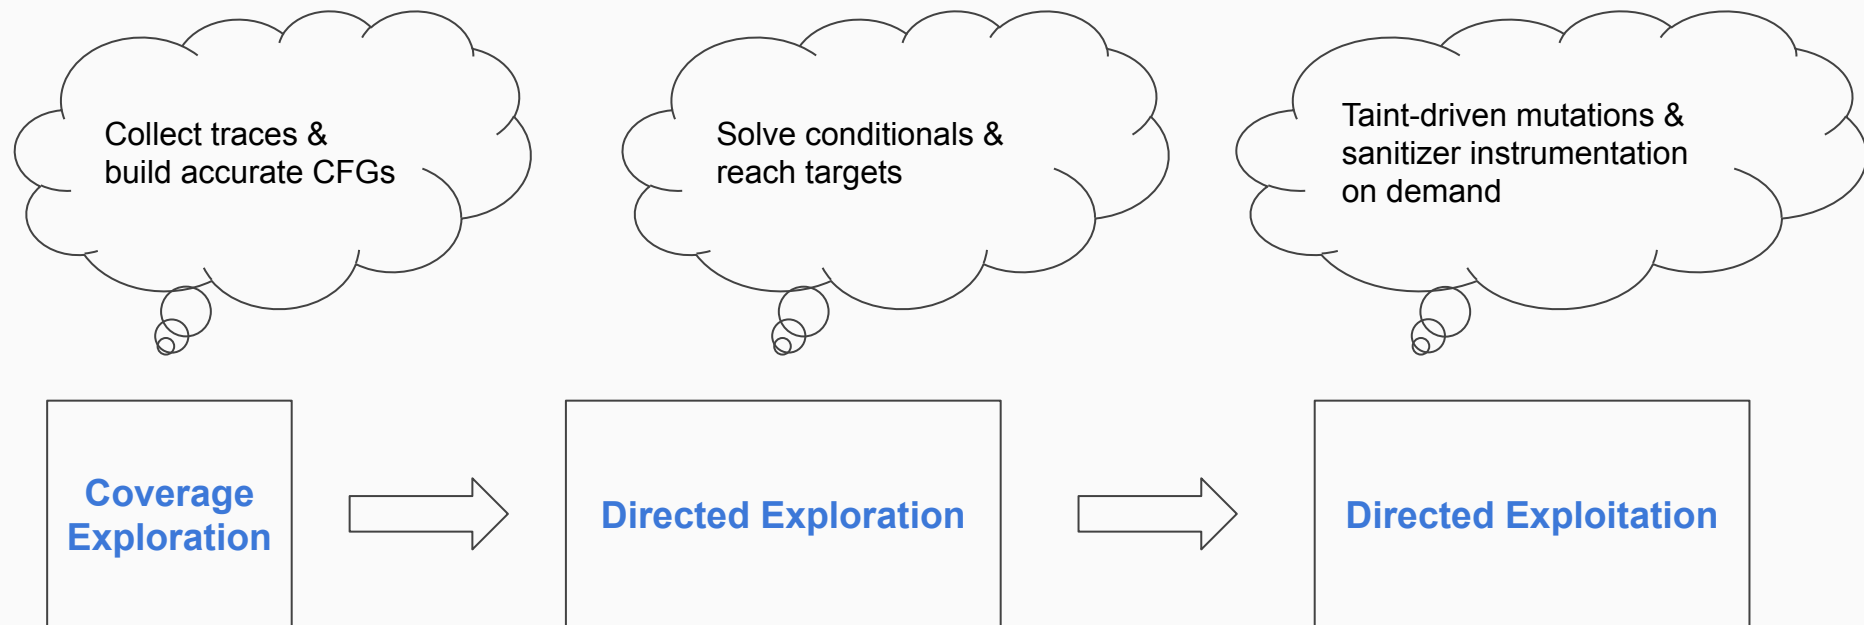


- SoA directed fuzzers rely on *statically-generated CFGs*
→ scale at runtime but *imprecise*

Goal: *Dynamically* construct *precise* CFGs

- Add edges *on the fly* → resolve indirect calls
- Construct Conditional Graph: compacted CFG that *only contains the conditionals* → use it to perform *scalable/precise* distance calculations
- Augment CFG with *taint analysis* → record a taint label at uncovered edges to decide mutate input bytes or not

Sanitizer-guided Fuzzer



Evaluation Configurations

- Research Questions
 - RQ1: [ParmeSan vs. Directed Fuzzers](#) (AFLGo, Hawkeye)
 - RQ2: [ParmeSan vs. Coverage-guided Fuzzers](#) (NEUZZ, QSYM, Angora)
 - RQ3: [Sanitizer Impact](#) (ASan, UBSan, LSan, TySan)
 - RQ4: [New Bugs](#)
- Implementation
 - ParmeSan = [LLVM passes](#) (static analysis) + [Angora](#) (fuzzer)
- Benchmark
 - [Binutils](#) (vs. directed fuzzers) + [Google Fuzzing Testsuite](#) (vs. coverage-guided fuzzers)

RQ1: ParmeSan vs. Directed Fuzzers

- Bench: 6 Binutils (AFLGo's paper) + 1 OpenSSL
- Rerun AFLGo
- Reclaim results of Hawkeye in paper

ParmeSan *outperforms directed fuzzers* against AFLGo's benchmark

CVE	Fuzzer	Runs	<i>p</i> -val	Mean TTE
OpenSSL				
2014-0160	ParmeSan	30	0.006	5m10s
	HawkEye	—		—
	AFLGo	30		20m15s
Binutils				
2016-4487 2016-4488	ParmeSan	30	0.005	35s
	HawkEye	20		2m57s
	AFLGo	30		6m20s
2016-4489	ParmeSan	30	0.03	1m5s
	HawkEye	20		3m26s
	AFLGo	30		2m54s
2016-4490	ParmeSan	30	0.01	55s
	HawkEye	20		1m43s
	AFLGo	30		1m24s
2016-4491	ParmeSan	10	0.003	1h10m
	HawkEye	9		5h12m
	AFLGo	5		6h21m
2016-4492 2016-4493	ParmeSan	30	0.003	2m10s
	HawkEye	20		7m57s
	AFLGo	20		8m40s
2016-6131	ParmeSan	10	0.04	1h10m
	HawkEye	9		4h49m
	AFLGo	5		5h50m

RQ2: ParmeSan vs. Coverage-guided Fuzzers

- Bench: 15 bugs of Google Fuzzing Testsuite
- All fuzzers run *with sanitizers enabled*
- AFLGo uses the targets obtained using the ParmeSan analysis stage

	Branch Cov	Time-to-Exposure
AFLGo	+16%	+288%
NEUZZ	+40%	+81%
QSYM	+95%	+867%
Angora	+33%	+37%
ParmeSan	-	-

- ParmeSan *outperforms coverage-guided fuzzers* (1) branch coverage (2) TTE
- Directed fuzzers require *less coverage*

RQ3: Sanitizer Impact

- Bench: 10 bugs
- ParmeSan using *different sanitizers* in the analysis stage.
- Example: UAF bug in pcre2
 - # targets: TySan < ASan (supports more bug classes)
 - TTE : TySan is 20% faster than ASan

Different sanitizers have *different impacts* on # targets and TTE

RQ4: New Bugs

- Bench: 12 targets from OSS-Fuzz to fuzz the most recent commits
- Found 47 unique bugs
 - 37 bugs in outdated library *pbc*
 - 10 bugs in well-fuzzed libraries

ParmeSan *found new bugs* in well-fuzzed programs

Limitations

- ParmeSan relies on LLVM IR and relevant analysis techniques
- Raw binaries: using binary hardening ?

- Composing sanitizers
- Apply heuristics in SoA directed fuzzers

Backup

Additional Results

Impact of different components

	Branch Cov	Time-to-Exposure
No lazysan	+0%	+25%
No pruning	+19%	+28%
No dyncfg	+17%	+34%
ParmeSan	-	-

Run-time and compile-time overhead

The overhead is negligible in most cases: *less than 3%* of the total execution time